

Inverse Reinforcement Learning Tool for Minigrid environment

Giulio Bazzanti Niccolò Biondi

Human Computer Interaction, February 2020

Table of Contents

- 1 Use Case
- 2 Proposed Method
- 3 IRL Tool
- 4 Experimental Results
- 5 Conclusions

Use Case
●○○○

Proposed Method
○○○○○○

IRL Tool
○○

Experimental Results
○○○○○

Conclusions
○○○

Use Case

Video games and AI

Reinforcement Learning is used in video games to create Non-Player Characters (NPCs).

Video games and AI

Reinforcement Learning is used in video games to create Non-Player Characters (NPCs).

Two problems:

- NPCs with super-human abilities
- NPCs with predictable behaviour

Video games and AI

Reinforcement Learning is used in video games to create Non-Player Characters (NPCs).

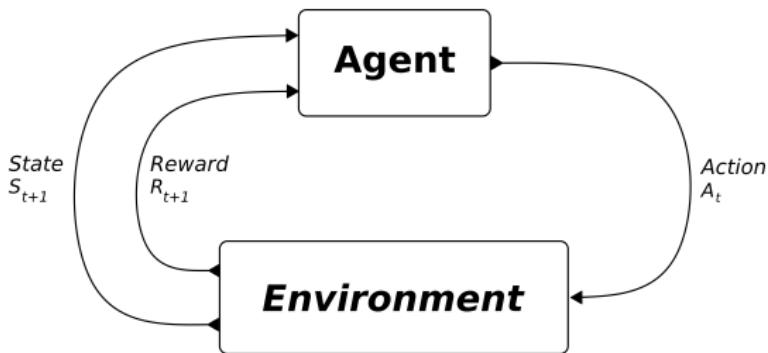
Two problems:

- NPCs with super-human abilities
- NPCs with predictable behaviour

Solution: Use in the agent train a human control

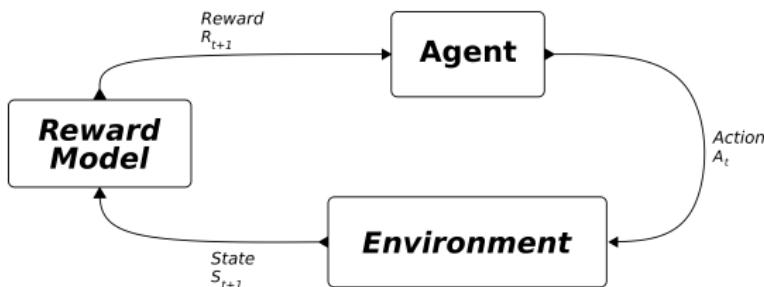
Reinforcement Learning

- Reinforcement Learning is the science of making optimal decisions using experiences
- Find an optimal behavior strategy for the agent to obtain optimal rewards



Inverse Reinforcement Learning

- The agent does not receive the reward from environment
- A human in the loop communicates to the agent his intentions



Proposed Method

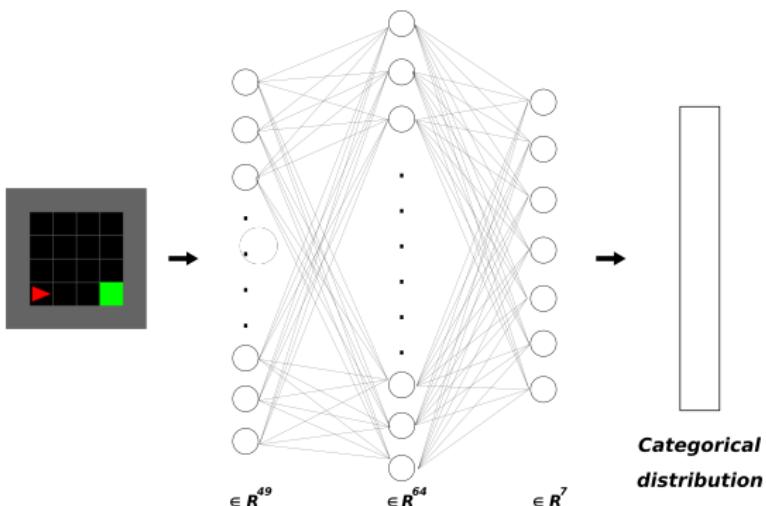
Proposed method

Algorithm 1 Training Protocol

- 1: Run the policy in the environment and store "initial trajectories".
 - 2: The annotator annotated all the "initial clips" and create the annotation buffer.
 - 3: Pretrain the reward model from the annotation buffer.
 - 4: **for** M epochs **do**
 - 5: Train the policy in the environment for N episodes with rewards from the reward model.
 - 6: Sample pairs of clips from the resulting trajectories.
 - 7: The annotator labels the selected pairs and puts them in the annotation buffer.
 - 8: Train the reward model for K batches from the annotation buffer.
 - 9: **end for**
-

Policy

- The reward is not taken from the environment but from a reward model
- The policy receives the current agent state as input and probabilities of actions as output



Annotator

- The annotator gives preference feedback about pairs of clips
 - (0,1) (1,0) preferred clips
 - (0.5,0.5) indifferent labels
 - (0,0) discarded clips

Annotator

- The annotator gives preference feedback about pairs of clips
 - (0,1) (1,0) preferred clips
 - (0.5,0.5) indifferent labels
 - (0,0) discarded clips
- Human annotator vs Artificial annotator (*Oracle*)

Annotator

- The annotator gives preference feedback about pairs of clips
 - (0,1) (1,0) preferred clips
 - (0.5,0.5) indifferent labels
 - (0,0) discarded clips
- Human annotator vs Artificial annotator (*Oracle*)
- All the preferences are stored in the Annotation Buffer

Reward Model

- The Reward Model has to emulate the annotator labels
- It is trained to minimize the cross-entropy loss between predictions and labels

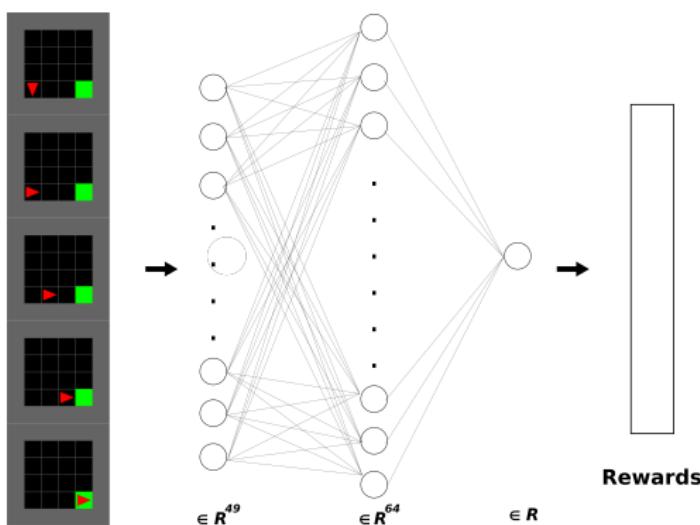
$$\text{loss}(\hat{r}) = - \sum_{(\sigma^1, \sigma^2, \mu) \in A} \mu(1) \log(\hat{P}[\sigma^1 \succ \sigma^2]) + \mu(2) \log(\hat{P}[\sigma^2 \succ \sigma^1])$$

- where

$$\hat{P}[\sigma^1 \succ \sigma^2] = \frac{\exp(\sum_{o \in \sigma^1} \hat{r}(o))}{\exp(\sum_{o \in \sigma^1} \hat{r}(o) + \sum_{o \in \sigma^2} \hat{r}(o))}$$

Reward Model

- The Reward Model predicts a list of rewards from a clip



IRL Tool

Use Case
oooo

Proposed Method
oooooooo

IRL Tool
○●

Experimental Results
ooooo

Conclusions
ooo

Bla Bla

Use Case
oooo

Proposed Method
oooooooo

IRL Tool
oo

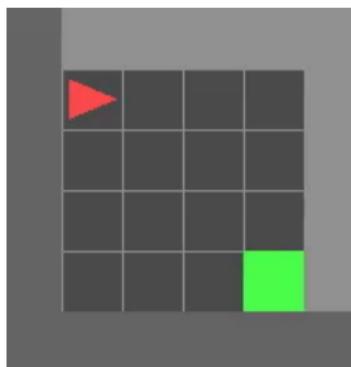
Experimental Results
●ooooo

Conclusions
ooo

Experimental Results

MiniGrid Environment

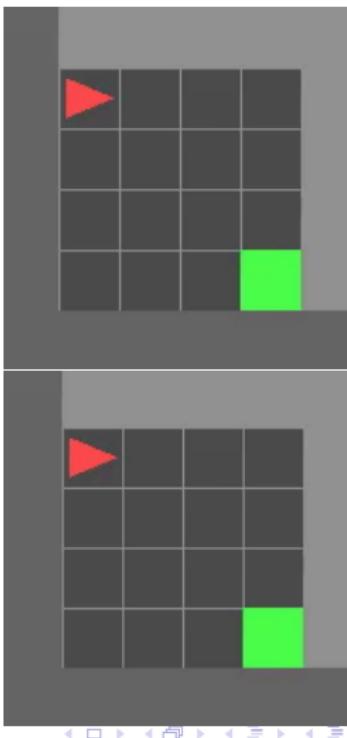
- 6x6 Empty MiniGrid environment for experiments
- Optimal trajectory achieved with RL



MiniGrid Environment

- 6x6 Empty MiniGrid environment for experiments
- Optimal trajectory achieved with RL

- Sub-optimal trajectory forced with IRL
- The user controls the agent behaviour



Tuning the components

- Policy

- the agent has to reach easily the environment goal \Rightarrow episode length 150
- no negative goal reward or big values range \Rightarrow standard deviation 0.5

Tuning the components

- Policy

- the agent has to reach easily the environment goal \Rightarrow episode length 150
- no negative goal reward or big values range \Rightarrow standard deviation 0.5

- Annotator

- sparse vs dense Oracle

Tuning the components

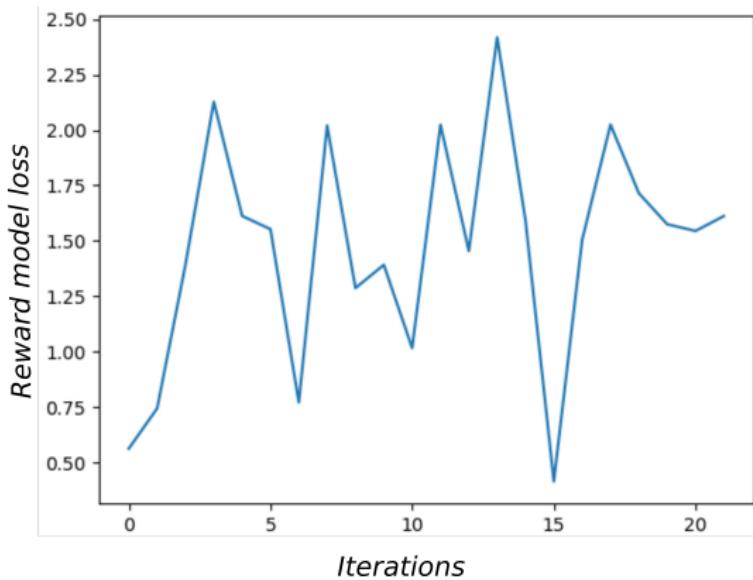
- Policy
 - the agent has to reach easily the environment goal \Rightarrow episode length 150
 - no negative goal reward or big values range \Rightarrow standard deviation 0.5

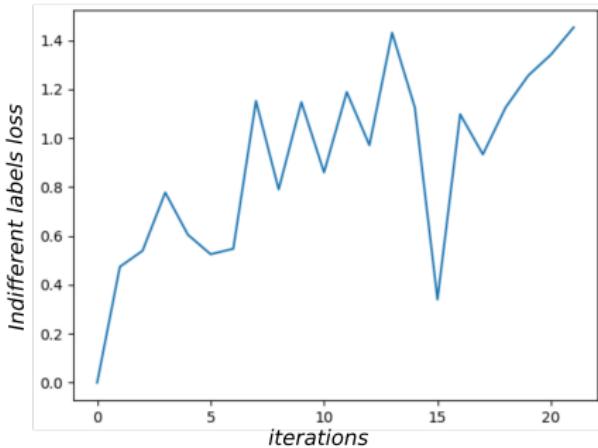
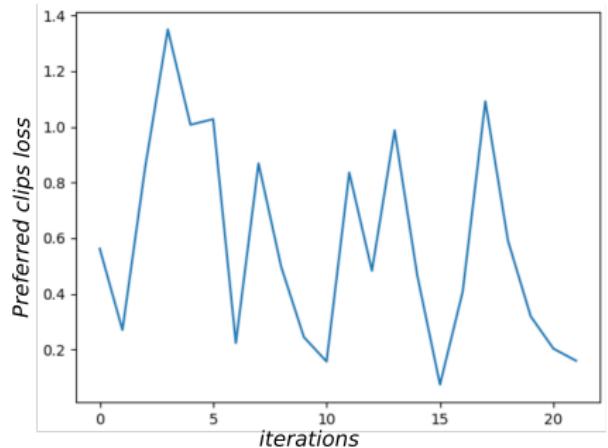
- Annotator
 - sparse vs dense Oracle

- Reward Model
 - variable K batch
 - early stopping

Reward Model Loss

- The reward model loss grows during the training protocol.





- Small contribution to total loss
- Same range during all the training

- Big contribution to total loss
- The number of those preferences make the loss grow

Conclusions

Conclusions

as

Future Works

cambiare ambiente (da cmabaire parametri e modelli? anche no dato che uno processa gli stati)
eventualmente processare anche le img in sequenza (cnn and lstm)